

Electronic warfare | **SIH26055**

Electronic warfare (**EW**) uses the electromagnetic spectrum to control space, disrupt enemy communications and radar, and protect friendly military forces.

Three Main Parts of Electronic Warfare:

- Electronic Attack (**EA**): Uses signals, jamming, or spoofing to block, confuse, or damage an enemy's radar, radios, and weapons guidance.
- Electronic Protection (**EP**): Shields your own team's communications and sensors from enemy interference or tracking.
- Electronic Support (**ES**): Listens to the airwaves to find, identify, and map out enemy signals and threats.

Why It Matters Today:

Modern weapons, drones, and navigation tools rely heavily on invisible radio waves and GPS. Controlling this spectrum is just as important as controlling the land, sea, or sky. If a military cannot protect its signals or block the enemy's signals, it can lose the battle before the fighting even starts.

Detailed Breakdown:

1. The Current Technology and Its Limit

- **The Routine:** Today's Electronic Support (ES) radar and radio receivers use a rigid, predetermined schedule to scan the skies for threats. This schedule is programmed into the system before the vehicle or aircraft ever starts its mission.
- **The Only Scenario Where It Works:** This static approach is completely fine if the enemy uses predictable, old, or unchanged radar equipment ("static, well-characterized threat environment").

2. Why Current Technology Fails in Modern War

The statement notes that pre-programmed scanning completely "falls apart" against advanced modern tactics because enemy emitters (radars and signal transmitters) do not cooperate. They disrupt fixed schedules using three tactics:

- **Frequency Hopping:** The enemy constantly changes their radio frequencies rapidly to dodge detection.
- **Intermittent Radiation:** They blink their radar systems on and off quickly, meaning a timed sweep might completely miss them when they are silent.
- **Intelligence Mismatch:** The enemy deploys new tech or operates in ways that completely contradict the older intelligence data loaded into the system before takeoff.

3. The Ultimate Goal ("The Ask")

The core objective of the Problem Statement (PS) is to move away from rigid automation and shift toward **cognitive electronic warfare**. Instead of following a fixed script, the software needs a smart brain that can autonomously **watch** the live spectrum, **learn** how the enemy is behaving on the fly, and **dynamically re-plan** its scanning schedule in real time.

Currently, the system is **rigidly automated** (it follows a strict script). The goal of the PS is to upgrade it to an **autonomous, cognitive system** (it thinks and adapts on its own).

The Automated System (What we have now)

An automated system is **fast but blind**. It operates strictly on rules written *before* the mission.

- **How it works:** It follows an exact, pre-programmed checklist. For example: *"Scan Frequency A for 2 seconds, then scan Frequency B for 2 seconds, then repeat."*
- **The Flaw:** It cannot adapt. If an enemy radar pops up on Frequency C, or only blinks when the receiver is looking at Frequency A, the automated system will miss it completely.
- **Human Dependency:** If the enemy changes tactics, a human programmer must manually rewrite the code or update the database before the next mission.

The Autonomous / Cognitive System (What the PS is asking for)

An autonomous, cognitive system uses Artificial Intelligence (AI) and Machine Learning (ML) to act like a **thinking digital operator**.

- **How it works:** Instead of a fixed schedule, it is given a *goal* (e.g., *"Find and track all active threats"*). It observes the environment, notices that an enemy radar is frequency hopping, and decides on its own to alter its scan schedule in microseconds to chase that signal.
- **The Advantage:** It learns and adapts *during* the mission. If it encounters a brand-new type of enemy signal that was never loaded into its memory database, it can analyze the pattern in real time and re-plan its strategy to counter it immediately.
- **Self-Sufficiency:** It does not need to wait for a human software update mid-battle.

Summary Comparison Table:

Feature	Rigidly Automated (Current Tech)	Autonomous / Cognitive (The Goal)
Control Logic	Strict pre-programmed script	Live AI / Machine Learning loops
Reaction to Surprise	Fails or ignores unknown threats	Analyzes, learns, and adapts
Scan Schedule	Fixed time intervals	Fluid and dynamic real-time changes
Best Used For	Predictable, old environments	Modern, unpredictable conflicts

The Core Truth (I)

An Electronic Support (ES) receiver is physically unable to listen to the entire sky at the same time. Scanning is a hard physical necessity, not a design choice.

1. The Physics Problem (The Bottleneck)

- **The "Eye" is Too Small:** Think of the radio spectrum as a massive highway, but the ES receiver's "eye" (Instantaneous Bandwidth) is only wide enough to look at one lane at a time.
- **The "Order of Magnitude" Gap:** In engineering, an "order of magnitude" means roughly 10 times. This means the slice of spectrum the receiver can look at *right now* is at least **10 times smaller** than the massive, full arena it is supposed to be protecting.

2. The Current Compromise (The Sweep)

Because it cannot see the whole picture at once, the receiver has to resort to **sweeping**.

- **The Process:** It tunes into Chunk 1, listens for a microsecond, hops to Chunk 2, listens, hops to Chunk 3, and so on.
- **The Gambling Strategy:** As the text states, it is essentially "**hoping**" that when it looks at Chunk 3, an enemy radar happens to be transmitting at that exact microsecond. If the enemy transmits while the receiver is looking elsewhere, the receiver is completely blind to it.

Connecting This to Previous Statement

- This is *why* the rigid schedule fails so badly.
- If the receiver follows a fixed, predictable script (e.g., always scanning lanes 1, 2, 3, 4 in order), an unpredictable enemy can simply time their transmissions to flash exactly when the receiver is looking away.
- By upgrading to a **cognitive system**, the receiver won't just "hope" and guess randomly; it will intelligently predict which lane the enemy is about to jump into and look there ahead of time.

The Core Truth (II)

This concept defines the architectural shift requested by the Problem Statement (PS). It moves the system from a dead-reckoning approach to an interactive, responsive loop.

1. Open-Loop System (The Current Status)

An open-loop system is **rigid, predictable, and non-responsive**.

- **The Process:** It takes an input (the pre-programmed plan), executes it, and outputs data. Crucially, the output never flows back to change the input.
- **The Downside ("Blindness"):** The text uses a great example: it treats a harmless background signal (like a civilian radio beacon) with the exact same priority as a lethal, weapon-guiding threat (fire-control radar). It wastes valuable time looking at empty air or useless noise because its instructions tell it to do so.
- **The Pro:** It requires very little computing power, making it incredibly fast and simple to run. (but blindly)

2. Closed-Loop System (The Goal)

A closed-loop system is **dynamic, prioritizing, and self-correcting**.

- **The Process:** It introduces a **feedback loop**. The system listens, detects a signal, and immediately uses that information to change its next move.

- **The Dynamic:** If it detects a tiny, suspicious blip on one frequency (a "hit"), it doesn't just move on to the next pre-planned chunk. It immediately rewires its own plan to stay on that frequency longer, or check back sooner, to verify if it is an incoming threat. Conversely, if a channel is completely empty ("misses"), it skips it to save time.
- **The Brain of the PS:** As the text states, this feedback loop is the absolute core mechanism you are being asked to build. Without a closed loop, cognitive electronic warfare cannot exist.

Summary Comparison Table

Feature	Open-Loop (Today's System)	Closed-Loop (The PS Target)
Feedback Mechanism	None. Ignores what it sees.	Continuous. Learn from hits/misses.
Time Management	Wastes time equally on safe and dangerous signals.	Prioritizes high-threat signals instantly.
Behavior	Predictable and static.	Adaptive and unpredictable to the enemy.
Analogy	Driving a car with your eyes closed, relying only on a map.	Driving with your eyes open, steering around obstacles.

The Core Truth (III)

Detecting an enemy signal is not a game of simple hide-and-seek; it is a game of **battleships played across two completely different grids at the same time: Space/Time and Frequency.**

If you are not at the exact right spot on both grids simultaneously, you are blind to the threat.

1. The Interception Matrix: Frequency x Time

To catch an enemy radar transmission, your receiver must achieve a perfect overlap in two dimensions:

- **The "Where" (Frequency):** Your receiver must be tuned to the exact radio channel the enemy is using.
- **The "When" (Time):** Your receiver must be listening to that channel at the exact millisecond the enemy decides to transmit.

If you match the frequency but look too late, you miss it. If you look at the right time but on the wrong channel, you miss it.

2. The Twin Countermeasures: How Enemies Exploit This

Modern radar systems are designed specifically to break this two-dimensional overlap, using two distinct tactics to keep the "where" and "when" constantly moving targets:

A. Frequency-Agile Emitters (Attacking the "Where")

- **The Tactic:** These radars do not stay on one radio channel. They "hop" across dozens of different frequencies every second.
- **The Impact:** Just as your receiver tunes into Channel A, the enemy leaps to Channel B. The "where" is a moving target, making a fixed, step-by-step sweep completely useless.

B. Spatially-Scanning Emitters (Attacking the "When")


- **The Tactic:** Even if an enemy radar stays on a single, predictable frequency, its physical or digital antenna is spinning or sweeping to search its own airspace.
- **The Impact:** The radar's energy behaves like a lighthouse beam. It only shines directly at your receiver for a tiny fraction of a second during its rotation. Even if you are tuned to the correct frequency, if you sweep past it while its "lighthouse beam" is pointed away from you, you detect nothing. The "when" becomes highly intermittent and unpredictable.

Connecting This to the Closed-Loop Goal

An open-loop, rigid system stands zero chance here. It sweeps blindly and hopes the enemy's frequency hops and antenna spins accidentally align with its pre-programmed schedule.

A **closed-loop, cognitive system** must track both dimensions. It needs to calculate the mathematical rhythm of the enemy's frequency hopping *and* the timing of its directional beam rotation. Once it figures out those two patterns, it can intelligently predict exactly **where** and **when** to look to catch the signal.

Here's roughly how the closed loop that the PS wants you to build actually fits together:



The Core Truth ()

The chart visually demonstrates the **closed-loop feedback architecture** required by the Problem Statement (PS). It shifts the receiver from a passive, blind listener into a data-driven, adaptive machine learning system. [1]

The metrics provided are the scoreboard. They prove mathematically that your intelligent AI scheduler can outperform a basic, mechanical sequential sweep (the "round-robin"). [2]

1. Visualizing the Closed-Loop Workflow

The diagram shows exactly how data flows dynamically to break the static routine:

- **RF Environment (The Threat):** The chaotic, live battlefield full of unknown, hidden enemy radar signals flashing across different bands.
- **ES Receiver (The Lens):** The hardware doing the physical work. Because its bandwidth is too narrow, it physically tunes into just **one band at a time**.
- **Hit / Miss Outcome (The Evaluator):** Did we catch a signal during that glance, or was it empty? This outcome is evaluated ("compared against ground truth") to measure real-time performance.
- **ML Scheduler (The Brain): This is your software target.** The machine learning brain absorbs the hit/miss data, calculates the pattern, updates its internal map, and determines the highest-priority band to look at next.
- **The Dashed Feedback Line (The "Ask"):** Instead of repeating a rigid, fixed cycle, the ML scheduler loops directly back into the ES receiver, dictating its next tuning move on the fly based entirely on what just happened.

2. The Success Metrics: Figures of Merit

To prove to a review board that your AI scheduler is better than a simple, dumb sequential sweep (round-robin), you must score your system using three categories:

A. Detection Quality (Is it accurate?)

- **Probability of Detection:** How many of the actual, real threats in the sky did your system successfully spot?
- **Probability of False Alarm:** How often did your system mistake harmless static or background noise for an enemy radar? (High false alarms waste vital system time).
- **Sensitivity:** How faint or weak can an enemy signal be before your system fails to notice it?

B. Speed and Efficiency (Is it fast?)

- **Average Intercept Time:** From the exact millisecond an enemy radar starts transmitting, how long does it take your system to tune in and lock onto it? (Lower is better).
- **Average Intercept Time Error:** How consistent is your speed? A system that is fast once but incredibly slow the next time has a high error rate and is unreliable.
- **Average Intercept Rate:** Over a sustained mission, what percentage of total enemy pulses/transmissions are you catching?

C. Learning Quality (Is it actually smart?)

- **Percentage of Correct Predictions:** How often did your ML model correctly guess where the frequency-agile or spatial-scanning emitter was going to jump next?
- **Average Reward / Cost:** A standard metric for Reinforcement Learning (RL). The system gets "points" (rewards) for catching critical threats quickly and suffers "penalties" (costs) for missing threats or wasting time on empty bands. High average reward means the AI is successfully optimizing its behavior.

The Ultimate Challenge: Defeating the "Round-Robin" Baseline

The judges evaluating your solution already know that a basic "round-robin" sweep (scanning Band 1, then 2, then 3, then 4, in a perfect circle) is very easy to build.

Your entire engineering argument depends on showing that your ML scheduler achieves **faster intercept times, higher detection probabilities, and lower time errors** against erratic, hopping emitters than a dumb sequential sweep ever could.

We have mapped out the core problem, the hardware constraints, the dimensions of the threat, and the architecture/success metrics of the solution.